




What Does an Anomaly Detector Add Beyond Simple Statistics?

Three distinct questions about the same detector score


Input and representations



Same series, two views


Q1 Standalone performance

Does the *detector* work?



Q2 Statistical overlap


How predictable is S_D from B ?



Q3 Conditional contribution

What does S_D add beyond B ?

$$CDU = L(B) - L(B, S_D)$$



our focus
matched held-out evaluation

Three questions

Main findings



Key takeaways